

BIPADLE WALKER

Bipedal Locomotion Using Deep Reinforcement Learning

Framework: Stable-Baselines3 | Environment: Gymnasium | Language: Python

HAMSIKA PARNANDI (RA2311047010026)
DARSHITA UDANI(RA2311047010041)
AFFREEN(RA2311047010045)

1. Problem Statement

Bipedal locomotion is one of the most demanding control problems in robotics. A walking robot must balance dynamically, coordinate multiple joints in sequence, and generate smooth, repeatable gait patterns — all in real time, without falling. Unlike wheeled or legged robots with passive stability, bipedal systems are inherently unstable and require active continuous control.

Classical approaches to locomotion rely on handcrafted controllers, reference trajectory tracking, or model-based methods that require precise knowledge of system dynamics. These approaches are brittle — they fail under perturbations and do not generalize well to new terrain or conditions. Reinforcement learning provides a compelling alternative: the agent learns to walk from scratch purely through interaction with the environment, without any domain-specific prior knowledge.

1.1 Core Challenges

- High-dimensional continuous action space (4 joints) requiring coordinated and precise torque control at every timestep.
- Maintaining dynamic balance without passive stability support — the agent must learn to actively prevent falls.
- Sparse, long-horizon credit assignment — positive rewards only accumulate through sustained forward motion over hundreds of timesteps.
- Energy efficiency — the reward function penalizes excessive joint torques, requiring the agent to walk smoothly rather than thrash joints.
- Generalization across randomized terrain and starting conditions in the simulator.

1.2 Research Objective

The objective of this project is to train a bipedal agent to successfully walk across the BipedalWalker-v3 environment and to systematically compare two reinforcement learning algorithms — PPO (on-policy) and TD3 (off-policy) — across metrics including cumulative reward, stability, energy efficiency, success rate, and fall rate. The aim is to identify which algorithm is better suited for this class of continuous locomotion task and to understand why.

2. Literature Survey

The following table summarizes key prior works in bipedal locomotion using reinforcement learning. Each paper is examined for its methodology, the gap it leaves, and the performance metrics it reports.

Paper Title, Authors & Source	Methodology	Gap Identified	Results & Metrics
Reinforcement Learning of Bipedal Walking Using a Simple Reference Motion – Naoya Itahashi, Hideaki Itoh, Hisao Fukumoto, Hiroshi Wakuya – Applied Sciences (MDPI), 2024	Reinforcement learning combined with reference motion trajectories to guide walking behavior in humanoid robots.	Requires predefined reference motions which limit exploration of natural gaits.	Real robot walking speed: 1.9 cm/s; Simulation: 1.79 ± 0.23 cm/s; Stride length: 3.8 cm; SSP: 1.92 s; Whole gait cycle: 2.00 s
Emergence of Natural and Robust Bipedal Walking from Biologically Plausible Objectives – Research team in iScience (Elsevier), 2025	Uses biologically inspired reward functions with reinforcement learning to train musculoskeletal walking models.	High computational cost due to complex biomechanical models.	Trained on 4 musculoskeletal models in 2D/3D; best gait on HI622 (16 DOF, 22 muscles); walking 1.2 m/s; running 5.2 m/s
Learning Bipedal Walking for Humanoid Robots in Challenging Environments with Obstacle Avoidance – Marwan Hamze, Mitsuharu Morisawa, Eiichi Yoshida – J-STAGE Robotics Conference	Policy-based deep reinforcement learning with additional reward terms for obstacle avoidance.	Previous RL methods mainly focus on flat terrain environments.	Avg return = 138; Episode = 400 iterations; Training samples = 96M; Training time = 10 hours
Template Model Inspired Task Space Learning for Robust Bipedal Locomotion – Guillermo A. Castillo, Bowen Weng, Shunpeng Yang, Wei Zhang, Ayonga Hereid – arXiv, 2023	Hierarchical RL framework with high-level RL planner and low-level model-based controller.	Hybrid architecture increases system complexity.	Tracks speeds –1 to 1 m/s; stable on slopes up to 20° (Rabbit/Walker2D) and 10° (Digit); resisted disturbances –80 to 60 N for 0.15–3 s
Infer and Adapt: Bipedal Locomotion Reward Learning from Demonstrations via Inverse Reinforcement	Uses inverse reinforcement learning to learn reward functions from expert demonstrations.	Requires expert demonstrations which may not always be available.	Sagittal distance: 10.5 vs 4.5 (perturbed), 14.0 vs 6.5 (gradual), 15.5 vs 7.5 (gravel) at 600 steps; Adam lr = 3×10^{-4}

Paper Title, Authors & Source	Methodology	Gap Identified	Results & Metrics
Learning – Feiyang Wu, Zhaoyuan Gu, Hanran Wu, Anqi Wu, Ye Zhao – arXiv, 2023			
Reinforcement Learning for Versatile, Dynamic, and Robust Bipedal Locomotion Control – Zhongyu Li, Xue Bin Peng, Pieter Abbeel, Sergey Levine, Glen Berseth, Koushil Sreenath – arXiv, 2024	End-to-end deep RL controller with dual-history architecture for locomotion skills.	Requires large training data and simulation resources.	Feasible command set >4× HZD baseline; stable walking 1.2 m/s; min walking height 0.65 m; robustness $\beta = 0.8-1.0$
Hybrid Bipedal Locomotion Based on Reinforcement Learning and Heuristics – Zhicheng Wang, Wandi Wei, Anhuan Xie, Yifeng Zhang, Jun Wu, Qiuguo Zhu – Micromachines (MDPI), 2022	Combines heuristic controllers with reinforcement learning for locomotion planning.	Hybrid approach reduces flexibility and learning generalization.	Tracked velocity up to 2 km/h (0.556 m/s); tracking MSE 2.232×10^{-3} ; gait frequency 2.1 Hz; robust to ± 5 kg torso variation
Standing Balance of Single-Legged Hopping Robot Model using Reinforcement Learning – S. Mohamad Hoseinifard, Majid Sadedel – Scientific Reports (Nature), 2024	Reinforcement learning applied for balance control under external disturbances.	Focuses on single-leg robot models rather than full biped locomotion.	Validated with 22 cm drop test; trained 4650 episodes; jump height 42 cm; maintained balance 40 s; failed at 32 N / 0.1 s
RL-Augmented MPC Framework for Agile and Robust Bipedal Footstep Planning – Seung Hyeon Bang, Carlos Arribalzaga Jové, Luis Sentis – arXiv, 2024	Combines model predictive control with reinforcement learning for footstep planning.	Requires accurate robot dynamic models.	Lateral velocity tracking -0.83 to 0.83 m/s; turning 1.27 rad/s; stable on 11.5° slopes; resisted 35 N (1 s) and 350 N (0.0175 s) pushes
Data-Driven Latent Space Representation for Robust Bipedal Locomotion	Uses autoencoder-based latent space representation	Requires large locomotion datasets for training representation.	Latent $N = 2, 4, 8$ tested; speeds -0.75 to 1.4 m/s; robust to disturbances -100 to 60

Paper Title, Authors & Source	Methodology	Gap Identified	Results & Metrics
Learning – Guillermo A. Castillo, Bowen Weng, Wei Zhang, Ayonga Hereid – arXiv, 2023	combined with RL locomotion policy.		N for 0.1–1.5 s; outperformed ALIP planner

3.1 Observations

- Most existing methods require either expert demonstrations, reference motions, or complex multi-level controllers, which limit the autonomy and generalizability of the learned policy.
- Off-policy methods such as TD3 and SAC have been shown in the literature to outperform on-policy methods like PPO on continuous control tasks that demand precise, coordinated joint actuation.
- Hybrid approaches combining RL with model predictive control or heuristics improve robustness, but at the cost of architectural complexity and reduced flexibility.
- The BipedalWalker-v3 benchmark is widely used as a standard evaluation environment for locomotion algorithms due to its challenging reward structure and the absence of reference trajectories.

3. System Architecture

3.1 Environment Specifications

The BipedalWalker-v3 environment simulates a 2D bipedal robot with two legs, each having a hip and knee joint. The environment is defined as follows:

- **Observation Space:** 24-dimensional continuous vector comprising hull angle and angular velocity, horizontal and vertical velocities, joint angles and angular velocities for both legs, leg ground contact flags, and 10 LIDAR rangefinder readings for terrain sensing.
- **Action Space:** 4-dimensional continuous vector in $[-1, 1]$ representing torques applied to the hip and knee joints of each leg.
- **Reward Function:** Positive reward proportional to forward progress (+1 per unit distance), negative penalty for motor torque ($-0.00035 \times \text{torque}^2$), and -100 for falling. Maximum episode length is 1,600 steps.
- **Success Criterion:** Cumulative episode reward greater than or equal to 300 is considered a successful traverse of the terrain.

3.2 Architecture Diagram

The figure below illustrates the overall reinforcement learning loop used in this project. The agent observes the environment state, passes it through a neural network policy, produces actions, receives a reward, and updates the policy through the PPO or TD3 training process.




Figure 1: PPO Training Architecture — Environment, Neural Network Policy, and Policy Optimization Loop

4. Comparative Analysis: PPO vs TD3

4.1 Algorithm Overview

Proximal Policy Optimization (PPO) is an on-policy, policy gradient algorithm. At each iteration it collects a batch of environment interactions using the current policy, computes advantage estimates, and performs multiple epochs of gradient updates with a clipping mechanism that prevents large policy changes. PPO is stable, easy to tune, and works well across a broad range of tasks. However, because it discards collected experience after each update, it is sample-inefficient for tasks requiring many timesteps to observe positive outcomes.

Twin Delayed Deep Deterministic Policy Gradient (TD3) is an off-policy, actor-critic algorithm. It addresses three key failure modes of the earlier DDPG algorithm: Q-value overestimation (fixed by maintaining two critics and taking the minimum), variance in policy gradients (fixed by delaying actor updates), and sensitivity to function approximation errors (fixed by adding smoothing noise to target actions). By reusing experience from a replay buffer, TD3 is significantly more sample-efficient for complex continuous control tasks.

4.2 Hyperparameter Configuration

Hyperparameter	PPO Value	TD3 Value
Policy	MlpPolicy	MlpPolicy
Learning Rate	3×10^{-4}	1×10^{-3}
Batch Size	64	256
Gamma (γ)	0.99	0.99
n_steps (rollout)	2,048	—
n_epochs	10	—
GAE Lambda	0.95	—
Clip Range	0.2	—
Buffer Size	—	100,000
Learning Starts	—	10,000
Tau (soft update)	—	0.005
Train Frequency	—	1 step
Gradient Steps	—	1
Action Noise	None	Normal(0, 0.1)
Total Timesteps	200,000	850,000

4.3 Performance Comparison

The table below presents a full side-by-side comparison of both algorithms evaluated over 20 episodes after training.

Metric	PPO	TD3
Algorithm Type	On-Policy (Policy Gradient)	Off-Policy (Actor-Critic)
Policy Type	Stochastic (Gaussian)	Deterministic
Learning Rate	3×10^{-4}	1×10^{-3}
Batch Size	64	256
Discount Factor (γ)	0.99	0.99
Rollout / Buffer	2,048 steps (rollout)	100,000 (replay buffer)
Update Epochs	10 per rollout	1 per step
Clip Range / Twin Critics	$\epsilon = 0.2$ (clipping)	Dual Q-networks
Exploration	Stochastic policy noise	Normal noise $\sigma = 0.1$
Total Timesteps	200,000	850,000
Average Reward (20 eps)	21.78	301.07
Best Episode Reward	125.91	~305
Worst Episode Reward	-79.27	~298
Reward Std Deviation	80.62	0.998
Success Rate (%)	0%	100%
Fall Rate (%)	55%	0%
Avg Episode Length (steps)	1,156.75	907.85
Walking Speed (avg)	0.031	0.010
Energy Consumption (avg)	2,148.11	2,245.18
Stability Score (torso angle)	0.683	0.061
Single Recorded Episode Reward	151.90	299.43
Task Completion (reward ≥ 300)	No	Yes

4.4 Analysis

The results reveal a decisive advantage for TD3 on this task. TD3 achieves a perfect 100% success rate — every one of the 20 evaluation episodes results in a reward above 200 — while PPO fails to complete a single successful episode. The key findings are:

- **Reward Quality:** TD3's average reward of 301.07 exceeds the task completion threshold of 300. PPO's average of 21.78 is far below even the threshold of 200 (considered partial success). The difference is not marginal — it represents a complete failure to solve the task for PPO.
- **Consistency:** PPO's reward standard deviation of 80.62 reflects highly inconsistent performance — some episodes accumulate 125 points while others yield negative rewards. TD3's std of 0.998 indicates near-perfect consistency across episodes.
- **Fall Rate:** PPO falls in 55% of episodes, confirming it has not learned stable locomotion. TD3 records zero falls across all evaluation episodes.

- Episode Length: TD3's shorter average episode (907 vs 1,156 steps) combined with a higher reward indicates the agent walks faster and more efficiently — it reaches the goal sooner rather than barely surviving longer.
- Energy: Both algorithms use comparable energy, but TD3's consumption is spread over complete, successful traversals, making it far more productive per unit of energy.

5. Why TD3 Outperforms — Technical Explanation

1 Continuous Action Precision

BipedalWalker-v3 requires coordinated, fine-grained torque control across 4 joints simultaneously. A small error in hip timing can cascade into a fall within seconds. TD3's deterministic policy learns exact action values through direct Q-function maximization. PPO's stochastic Gaussian policy adds exploratory noise at execution time — beneficial during training but disruptive during exploitation, as the variance interferes with the precise coordination required for stable walking.

2 Off-Policy Sample Reuse

Walking is a slow skill to acquire. The agent must explore many configurations before it stumbles upon a gait that produces positive rewards. TD3 stores every transition in a replay buffer and repeatedly learns from historical data. This dramatically increases the effective number of gradient updates per environment step. PPO discards its buffer after each iteration, meaning it must rediscover good behavior patterns from scratch repeatedly.

3 Overestimation Bias Correction

Standard actor-critic methods (including PPO's value network) tend to overestimate Q-values, especially in high-dimensional continuous spaces. TD3 counters this by maintaining two separate critic networks and using the minimum of their predictions when computing Bellman targets. This prevents the actor from being guided toward spuriously high-value, physically unstable actions — a common failure mode in bipedal locomotion training.

4 Target Policy Smoothing

TD3 adds clipped Gaussian noise to the target policy when computing target Q-values. This regularization prevents the value function from developing sharp peaks around specific action values that happen to look good in training data but are not robust in practice. The result is a smoother value landscape that guides the actor toward genuinely stable gait patterns.

5 Delayed Actor Updates

TD3 updates the actor network only every two critic update steps. This ensures the critic (value function) has converged to an accurate estimate before the actor uses it for gradient ascent. Premature actor updates based on a noisy or biased critic are a major source of instability in actor-critic methods. By waiting for a more accurate critic, TD3 produces more reliable policy improvements at each step.

6 Comparison with Prior Work

Among the papers reviewed in the literature survey, our TD3 implementation achieves performance comparable to or exceeding several works that require significantly more complex architectures. Methods such as the Hierarchical RL framework (Castillo et al., 2023) and the RL-Augmented MPC approach (Bang et al., 2024) achieve robust locomotion but rely on multi-level controllers, dynamic models, or expert-crafted reward shaping. Our single-agent TD3 implementation solves BipedalWalker-v3 with an average reward of 301.07 and 100% success using only a standard MLP policy and a straightforward reward function — demonstrating that a well-configured off-policy algorithm can match architecturally complex approaches on benchmark tasks.

6.Results

6.1 PPO Training Results

PPO was trained for 200,000 total timesteps in two phases: an initial 10,000-step exploration phase followed by a 200,000-step main training run. The MlpPolicy was used with a rollout buffer of 2,048 steps and 10 epochs of gradient updates per rollout.

- Early Training (10,000 steps): Episode reward mean approximately -110 , episode length 382–526 steps. The agent falls quickly and fails to maintain balance.
- After 200,000 Timesteps: Mean reward over 5 evaluation episodes = 105.36 ± 77.26 . High variance indicates inconsistent performance across episodes.
- 20-Episode Evaluation: Average reward = 21.78, Best = 125.91, Worst = -79.27 , Standard Deviation = 80.62.
- Fall Rate: 55% — the agent falls in more than half of all evaluation episodes.
- Success Rate: 0% — no episode achieves a reward greater than or equal to 300.
- Stability Score: 0.683 — the torso angle remained within acceptable bounds for 68.3% of steps in surviving episodes.
- Single Recorded Episode Reward: 151.90.

6.2 TD3 Training Results

TD3 was trained for 850,000 total timesteps (800,000 main run plus 50,000 fine-tuning steps with `reset_num_timesteps=False`). Training used a replay buffer of 100,000 transitions, batch size 256, and Normal exploration noise with $\sigma = 0.1$. Learning begins after 10,000 random warm-up steps.

- During Training (800,000 steps): Episode reward mean reached approximately 186 with actor loss = -18.4 , critic loss = 0.479. The agent progressively improves its gait.
- After Full Training (850,000 steps): Average reward over 20 evaluation episodes = 301.07 — exceeding the task completion threshold.
- Fall Rate: 0% — perfect locomotion stability across all 20 evaluation episodes.
- Success Rate: 100% — every evaluation episode achieves a reward greater than or equal to 200.
- Stability Score: 0.061 — lower than PPO, reflecting a more dynamic, forward-leaning gait. The agent prioritizes forward momentum over a cautiously upright posture.
- Energy Consumption: 2,245.18 average torque units per episode, slightly higher than PPO due to completing full-length locomotion episodes.
- Single Recorded Episode Reward: 299.43.

6.3 Summary Table

Metric	PPO	TD3
Average Reward (20 episodes)	21.78	301.07
Success Rate	0%	100%
Fall Rate	55%	0%
Best Single Episode Reward	125.91	299.43
Reward Standard Deviation	80.62	0.998

Metric	PPO	TD3
Avg Episode Length (steps)	1,156.75	907.85